

# UCF ARCC Stokes Reference Card

March 19, 2017

Created and maintained by [Armando Fandango](#), available from [GitHub](#). Feel free to contact him for contributions.

## Getting help

ARCC helpdesk: [request-stokes@ist.ucf.edu](mailto:request-stokes@ist.ucf.edu)

## Stokes - Logging In

```
ssh -Y -i path_to_your_ssh_key your_username@stokes.ist.ucf.edu
    For Mac/Unix/Windows based system with ssh client
nano ~/.ssh/config For setting up personal SSH configuration
Host stokes
  Hostname stokes.ist.ucf.edu
  user your_username
  IdentityFile absolute_path_to_ssh_key
```

## Stokes - Job Queues

queue	Max Cores	Max Nodes	Max Time
batch64	512	64	240:00:00
batch98	784	98	48:00:00
batch224	2688	224	36:00:00

## SLURM - Job Submission

**sbatch** Submit script for later execution  
    --array=<indexes> Job Array Specification

**srun** Create a job allocation and launch a job step

    --label Prepend task ID to output

**salloc** Create job allocation and start a shell

**Options common across sbatch, srun and salloc:**

    --account=<name> Account to be charged for resources used  
    --begin=<time> Initiate job after specified time  
    --constraint=<features> Required node features  
    --cpu\_per\_task=<count> Number of CPU's required per task  
    --dependency=<state:jobid> defer until jobid reaching specified state  
    --error=<filename> File to store job errors  
    --exclude=<hostnames> Hosts to be excluded from job allocation  
    --export=<name[=value]> Export specified environment variable name  
    --gres=<name[:count]> Generic resources required for the job  
    --input=<name> File for input job data  
    --job-name=<name> A name for the job

## SLURM - Job Management

**sattach** Connect stdin/out/err for an existing job or step  
**scancel** Cancel jobs or steps

    --account=<name> Operate only on jobs for the specified account  
    --name=<name> Operate only on the specified job(s)  
    --partition=<name(s)> Operate only on the jobs in specified partition(s)  
    --qos=<name> Operate only on the jobs using specified QoS  
    --reservation=<name> Operate only on specified reservation  
    --state=<name(s)> Operate only on job in specified state  
    --user=<name> Operate only on jobs for the specified user  
    --odelist=<name(s)> Operate only on jobs using the specified node(s)

**sbcast** Transfer file to compute node(s) allocated to a job  
**srun\_cr** Wrapper for srun to support Berkeley checkpoint/restart  
**strigger** Event trigger management create, destroy or list event triggers

## SLURM - Get Information

**sinfo** System status - nodes, queues

**squeue** Job and Step status

    --account=<name> Operate only on jobs for the specified account  
    --jobs=<job id list> Operate only on the specified job(s)  
    --name=<job name(s)> Operate only on the specified job(s)  
    --partition=<name(s)> Operate only on the jobs in specified partition  
    --priority Sort jobs by priority  
    --qos=<name> Operate only on the jobs using specified QoS  
    --reservation=<name> Operate only on specified reservation  
    --state=<name(s)> Operate only on job in specified state(s)  
    --users=<name(s)> Operate only on jobs for the specified user(s)  
    --odelist=<name(s)> Operate only on jobs using the specified node(s)

**smap** Status with curses-based GUI

**sviiew** Status with GTK-based GUI

**sacct** Accounting info by job and step

**sstat** Accounting info about currently running jobs and steps

**sreport** Resource usage by cluster, partition, user, account etc.

## SLURM - Environment Variables

SLURM_ARRAY_JOB_ID	Job Id in Array
SLURM_ARRAY_TASK_ID	Task Id in Array
SLURM_CPUS_PER_TASK	No. of CPUs per task
SLURM_JOB_ACCOUNT	Account Name
SLURM_JOB_ID	Job Id
SLURM_JOB_NAME	Job Name
SLURM_JOB_NODELIST	Names of nodes allocated
SLURM_JOB_NUM_NODES	No. of nodes allocated
SLURM_JOB_PARTITION	Partition running the job
SLURM_JOB_UID	User Id of job owner
SLURM_JOB_USER	User name of job owner
SLURM_NPROCS	Number of cores allocated
SLURM_PROCID	Task Id (MPI Rank)
SLURM_STEP_ID	Step Id in a Job
SLURM_STEP_NUM_TASKS	No. of tasks (MPI Ranks)

## SLURM - Administration

**scontrol** Admin tool to view/update system, job, step, partition, reservation status

**sacctmgr** DB Admin tool to add/delete clusters, accounts, users and get/set resource limits, fair-share allocations etc.

**sprio** View factors comprising job's priority

**sshare** View current hierarchical fair-share information

**sdiag** View statistics about scheduling module

## Torque/MOAB - User Commands

myquota	Check your disk usage
mysusage	Check your hours usage
qsub	submit job
qstat   showq	status of queue
pbsnodes   checknode	status of node
checkjob	status of node
mdiag	

## StarCCM

Make sure the following lines are added to your ~/.ssh/config file:

```
Host *
StrictHostKeyChecking no
UserKnownHostsFile=/dev/null
```

## Matlab

Compile: Get an interactive node, load Matlab module and then compile the Matlab code.

Run: Submit the compiled code as a SLURM job.

**mcc -m hello.m** Compile matlab code as standalone C application

**mcc -p hello.m** Compile matlab code as standalone C++ application